

The Integration Game

How to Pick a Winning Integration Pattern



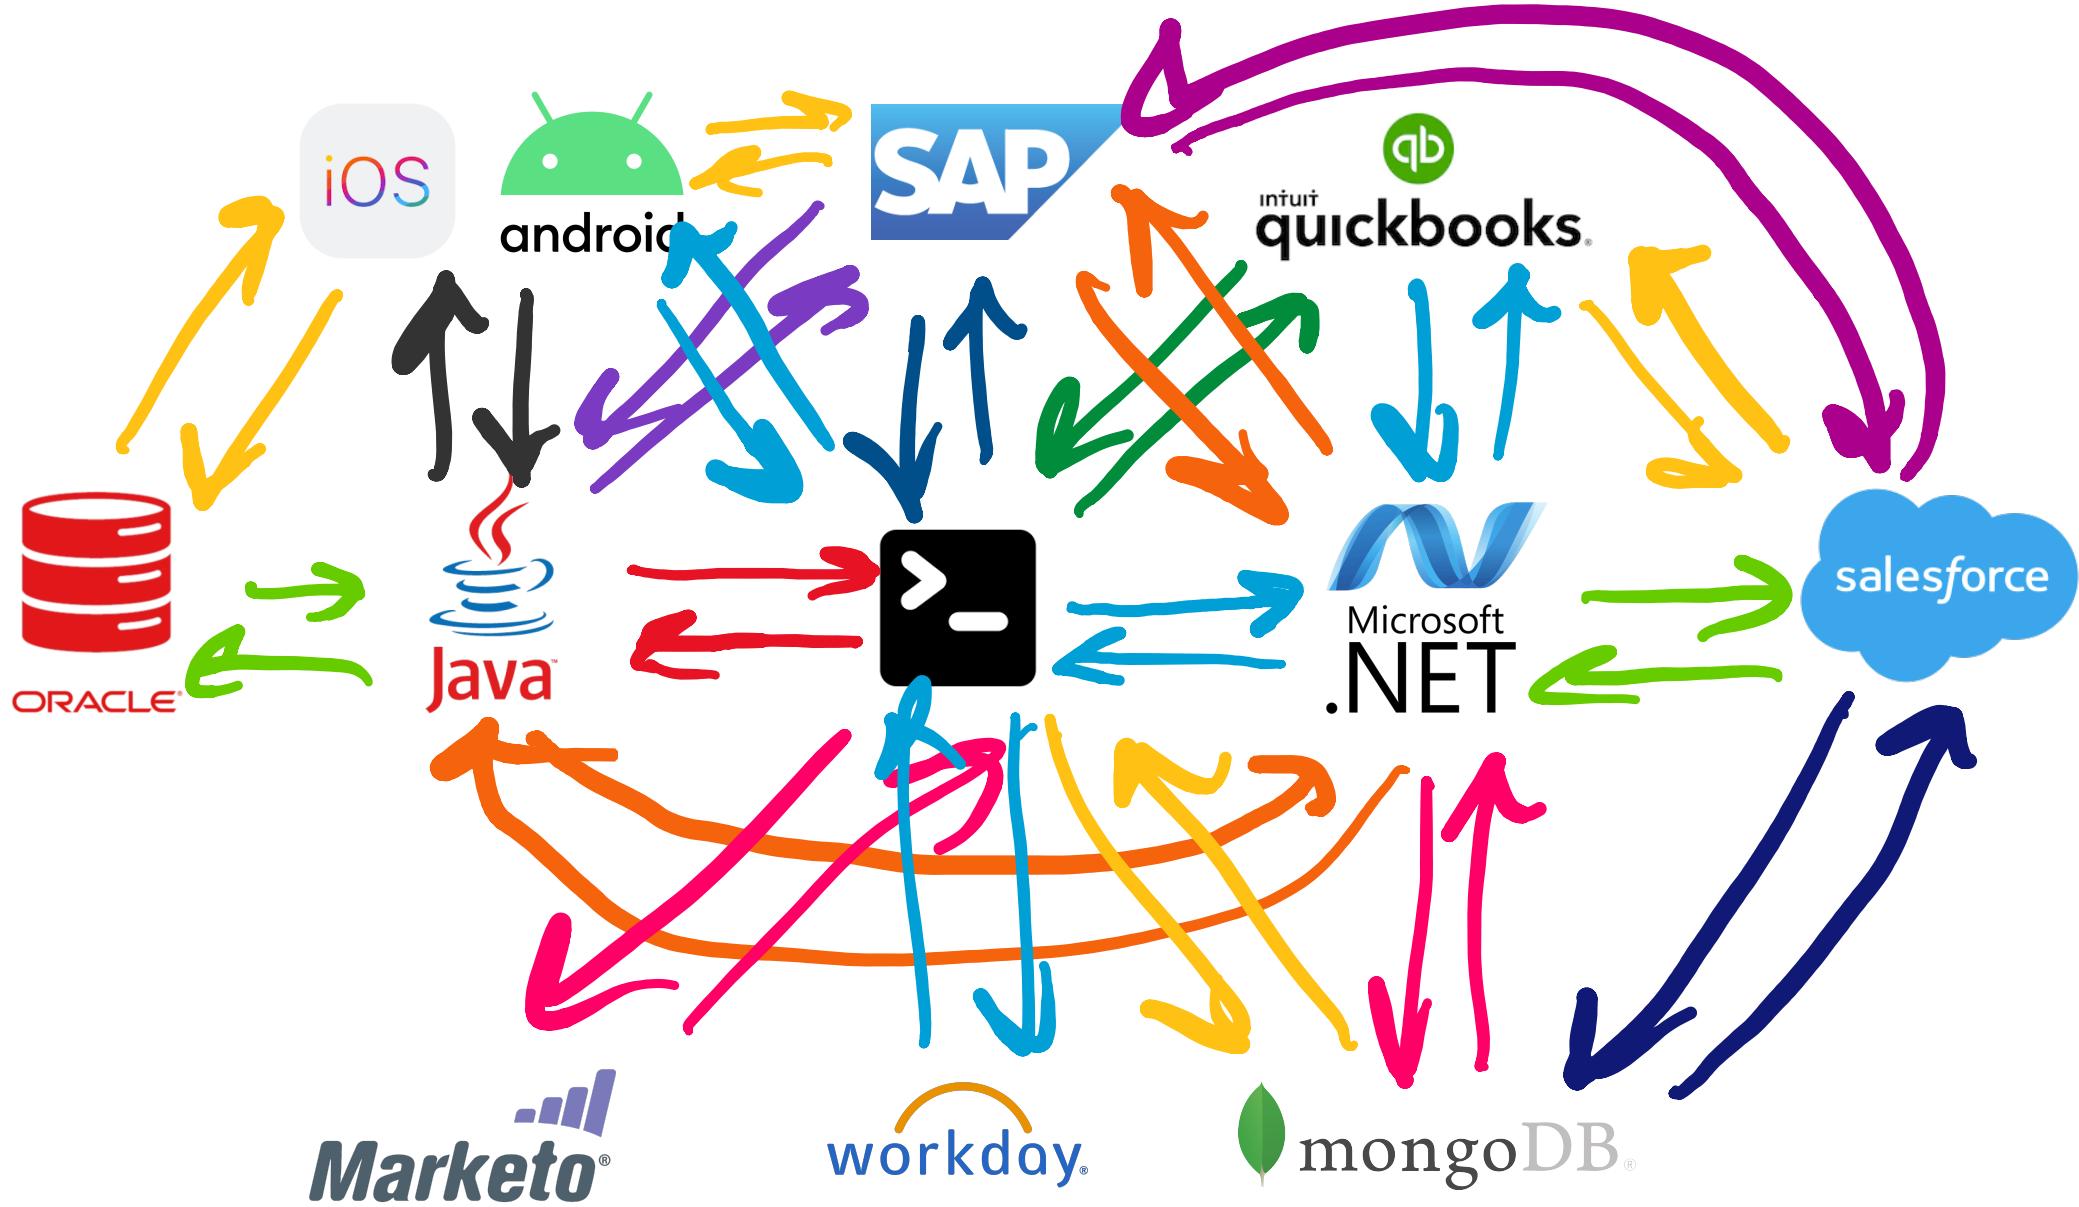
Branden Bellanca
Salesforce Architect



@bsbellanca



/branden-bellanca





"We'll fix it in Phase 2"



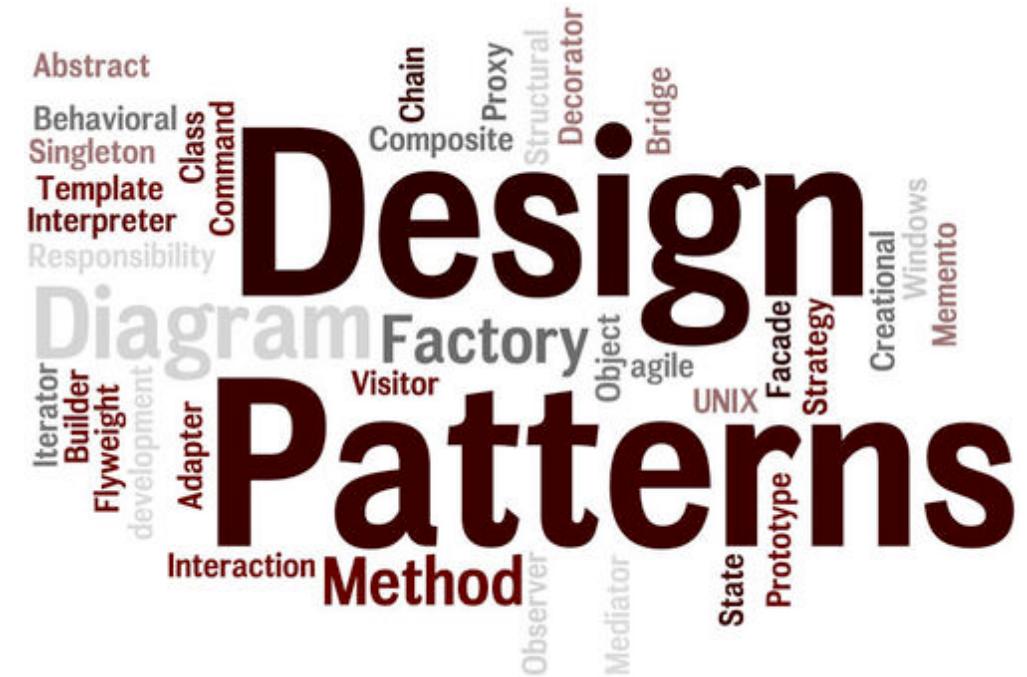
Makeshift Integrations

The "Great Value" Guide

@bsbellanca

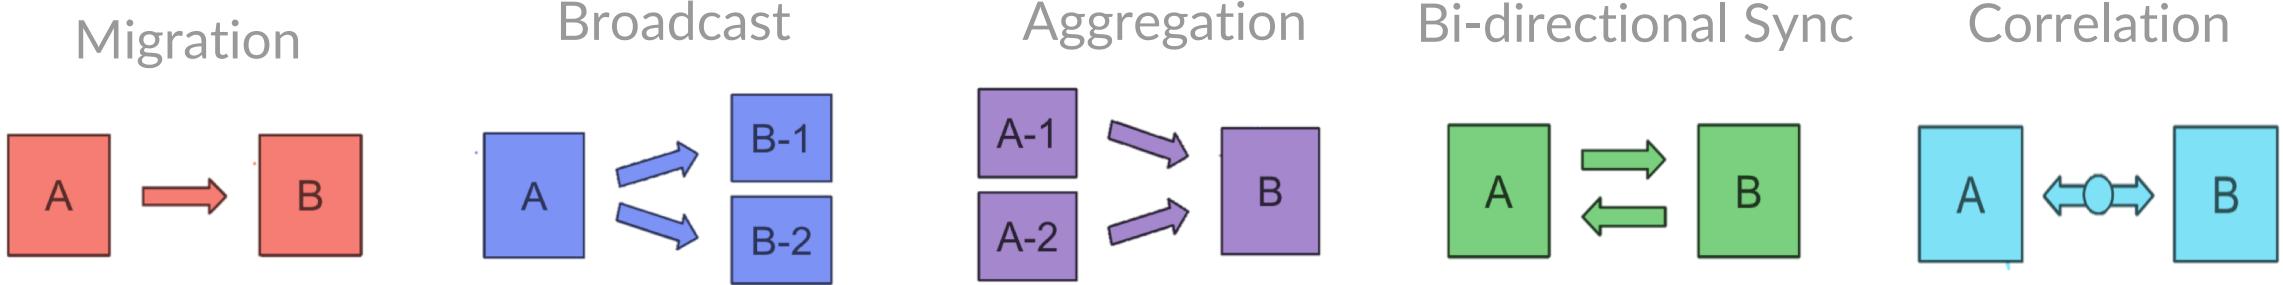
Design Patterns

- ✓ Represent some of the best practices adapted by experienced developers
- ✓ Systematically names, motivates, and explains a general design that addresses a recurring design problem
- ✓ Describes the problem, the solution, when to apply the solution, and its consequences



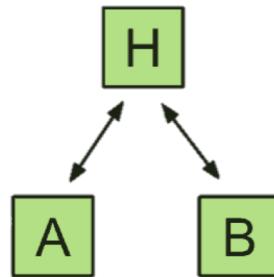


Integration Patterns

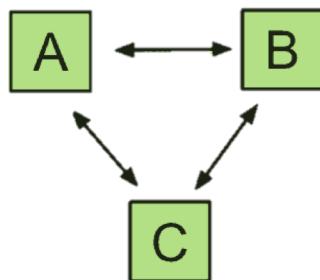


INTEGRATION ARCHITECTURE

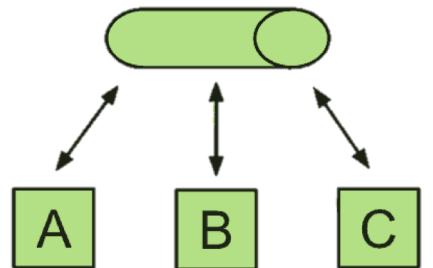
Point-to-Point



Hub and Spoke



ESB

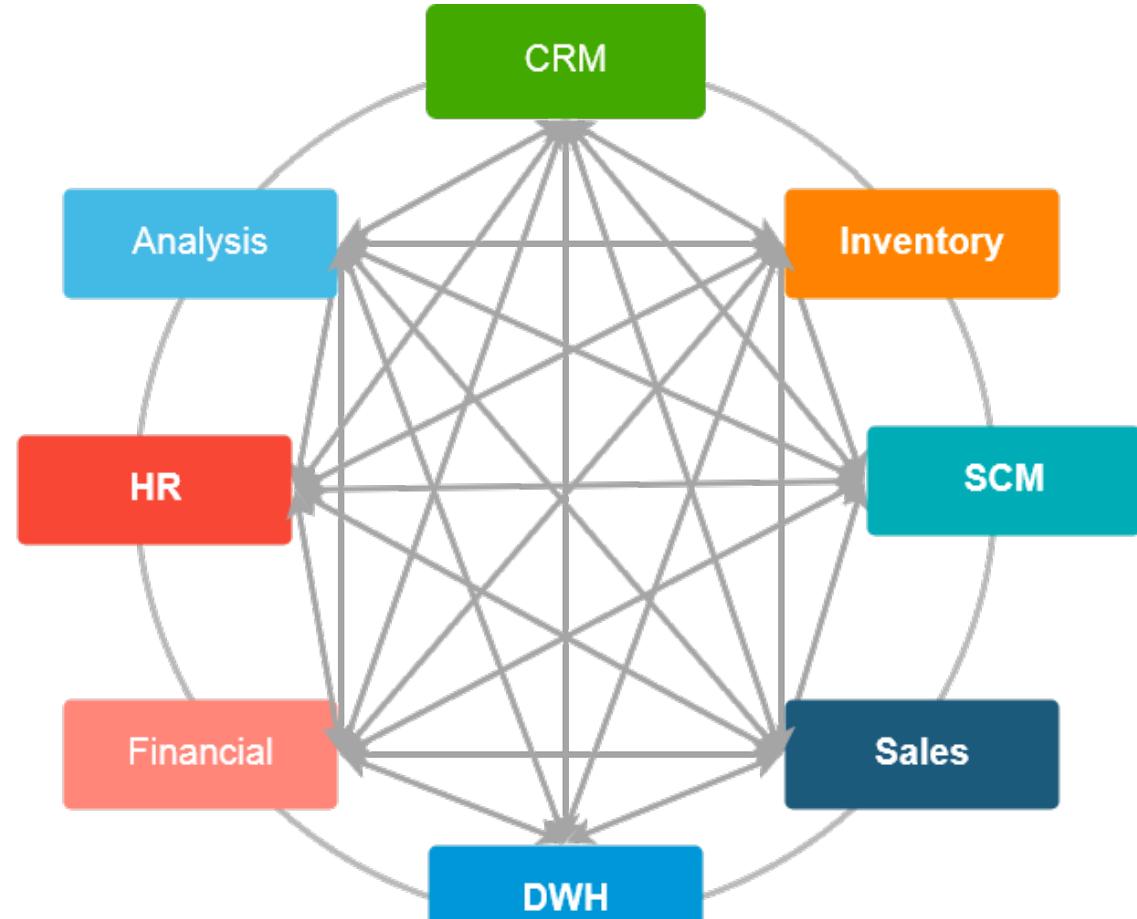




INTEGRATION ARCHITECTURE

Point-to-point

- ✓ Data flows directly from system to system
- ✓ Easy to implement
- ✓ Very hard to scale



Point-to-point
Integration

INTEGRATION ARCHITECTURE

Hub and Spoke

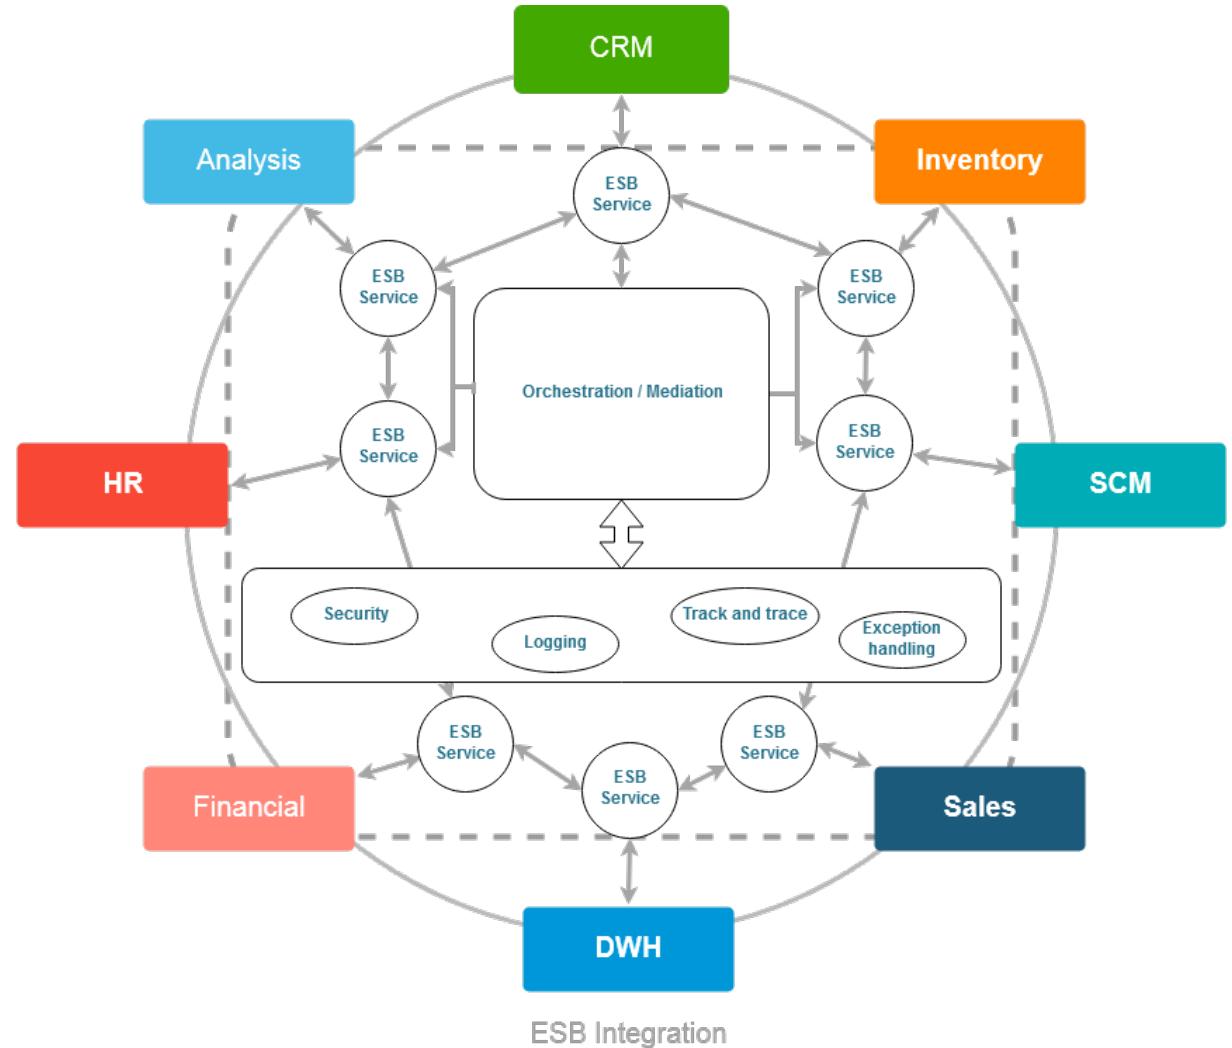
- ✓ Data flows through a central point
- ✓ Single point of failure
- ✓ Adds another place where development must take place and another runtime component



INTEGRATION ARCHITECTURE

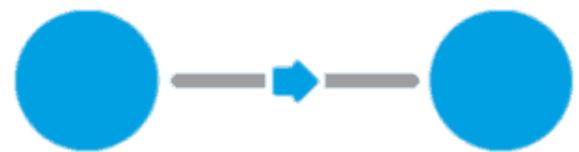
Enterprise Service Bus (ESB)

- ✓ Distributed services architecture
- ✓ All systems follow the same standards
- ✓ Any new system can plug into the bus
- ✓ Highly scalable



Migration

The act of moving a specific set of data at a point in time from one system to the other



What is it?

Source system
where the data
resides prior to
execution

1

Criteria which
determines the
scope of the data
being migrated

2

A transformation
that the data set
will go through

3

A destination
system where
the data will be
inserted

4

Ability to capture
results of the final
state vs the desired
state

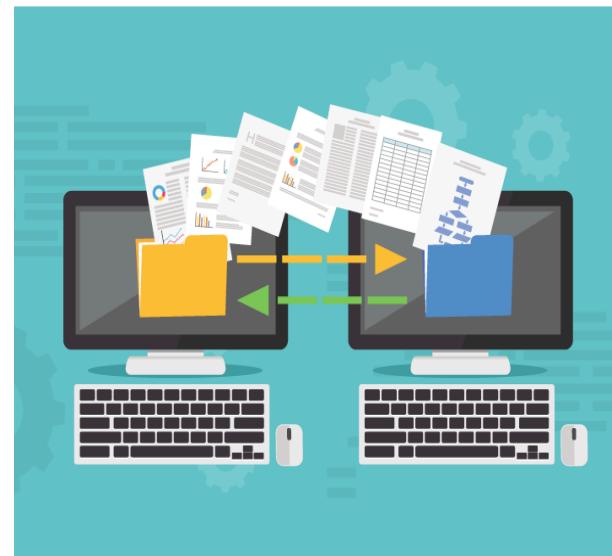
5

Some Clarification



Data Migration

the act of moving data
between systems



Application Migration

the act of moving
functionality capabilities
between systems

Business Data vs Business Capabilities

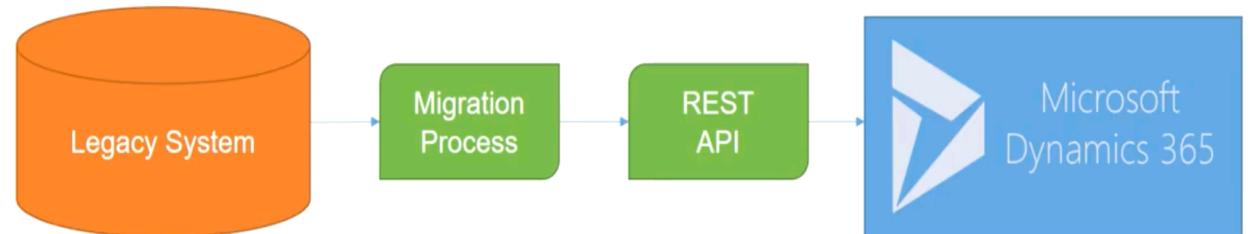
Is this really integration?

In the context of one-off, one-way copy of data from **one legacy system to another newer system**... The argument is unconvincing

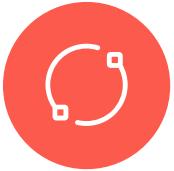


A special type of integration

- ✓ Migration is truly **moving** a set of data at a point in time (a snapshot) from one system to another
- ✓ This **move** may occur multiple times



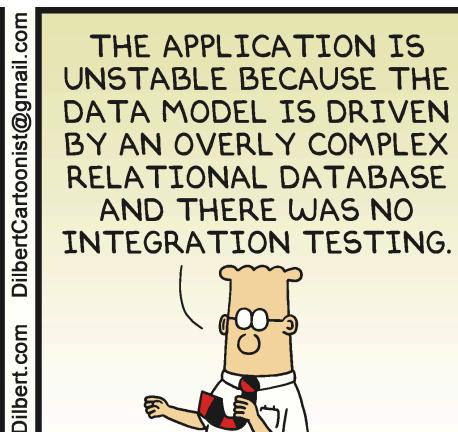
Why is it valuable?



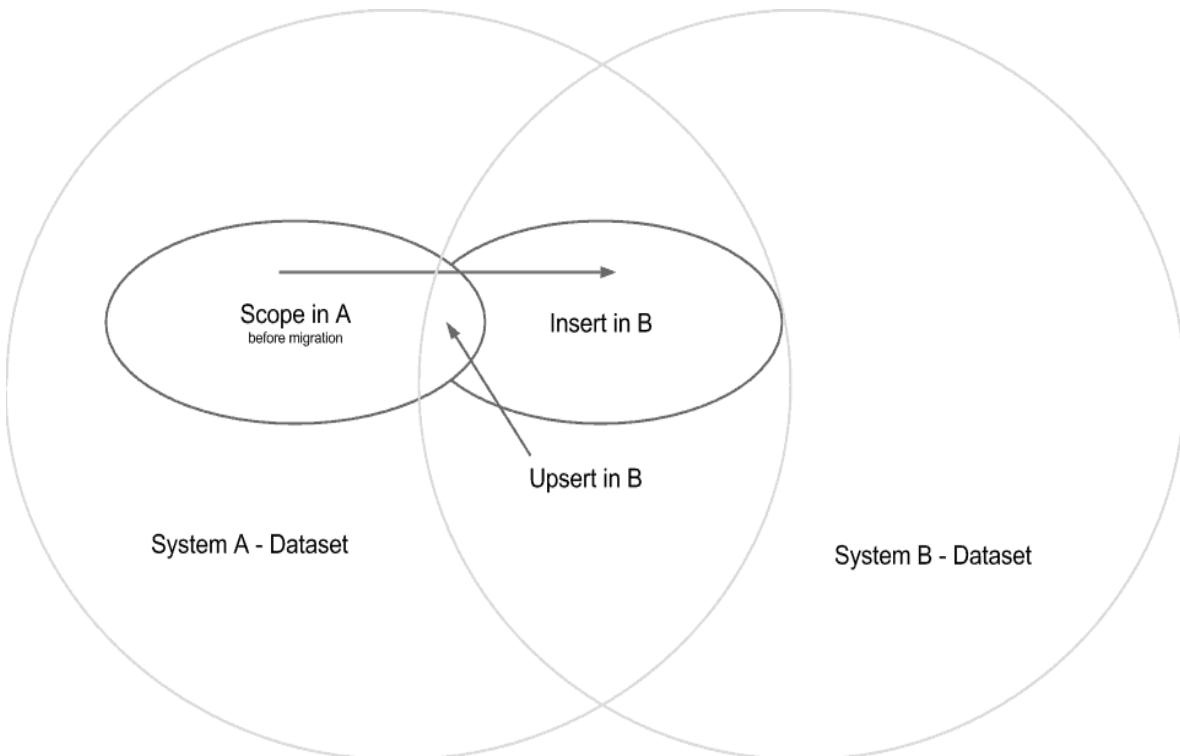
Keeps data agnostic from the tools that we use to create, view, and manage it



Without migration, we would lose any amassed data anytime that we changed systems



When is it useful?



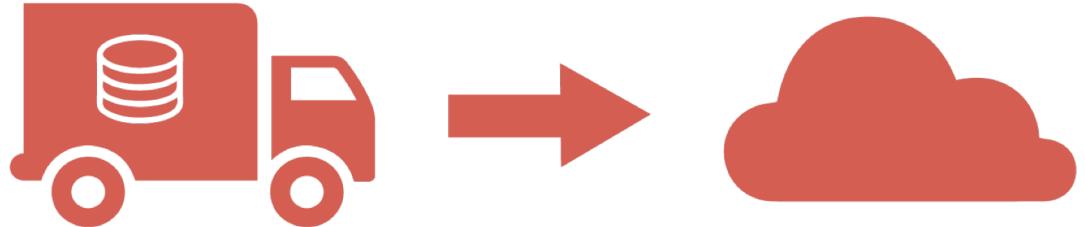
- **Moving from one system to another**
- **Moving from one instance of a system to another**
- **Backing up a dataset**
- **Adding nodes to database clusters**
- **Replacing database hardware**
- **Consolidating systems**

Consideration: Scoping the Dataset

✓ Migrations are a perfect time to clean up or at least move only the valuables to the new destination.

✓ It's tempting to just move everything, which tends to increase the scope of migrations in a futile effort to save potentially data with 0 value

✓ working with the data owner to reduce the scope of the data needing to be migrated can ensure only data that adds value is migrated



Consideration: Pull or Push?



Push Model

“once a specific event happens in the origin system, push all the data that matches a specific criteria to a processing app that will then insert it into the destination system”

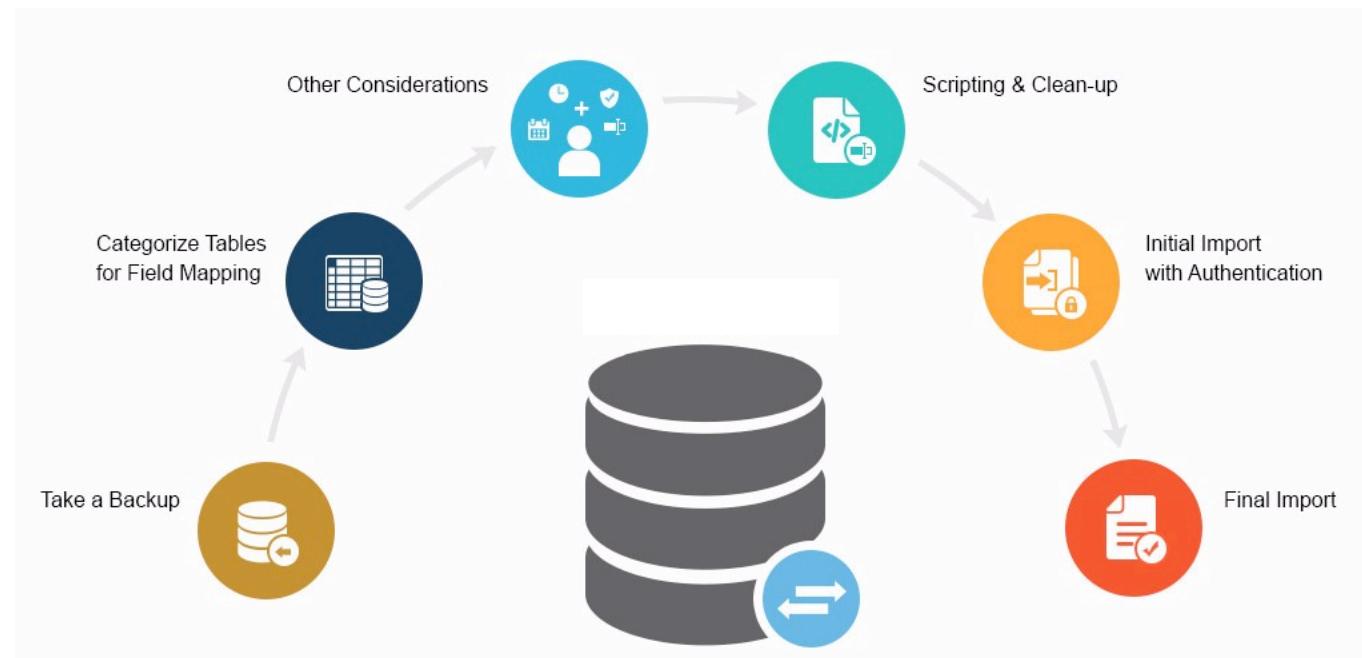
Pull Model

“let the data collect in the origin system until I am ready to move it, once I am ready, pull all the data that matches my criteria and move it to the destination system.”



Consideration: Triggering the Migration

- ✓ Traditionally handled via scripts, custom code, or database management tools, that developers or IT operations will create
- ✓ iPaaS systems, such as Mulesoft or Dell Boomi, can capture migration scripts or apps as integration apps
- ✓ Having services for common migrations save your team a lot of time



What is an iPaaS?

Integration Platform as a Service

A cloud integration platform as a service, enabling connectivity to SaaS and cloud services and providing a secure method of accessing on-premises applications behind a firewall

On-Premises



ERP/CRM



Oracle Data Warehouse



Home-Grown Apps



iPaaS

- DevOps
- Admin
- Integration Platform Services
- Governance Platform Services
- Cloud Foundation Services
- Collaboration, Self-Service, Integration Marketplace



Design Tool



Runtime Agent

Cloud



Salesforce



AWS Redshift



Open Data



Cost-efficiency

Reduces the need to generate code for customized integrations



Scalability

Can expand without having to worry about setting up another custom in-house integration

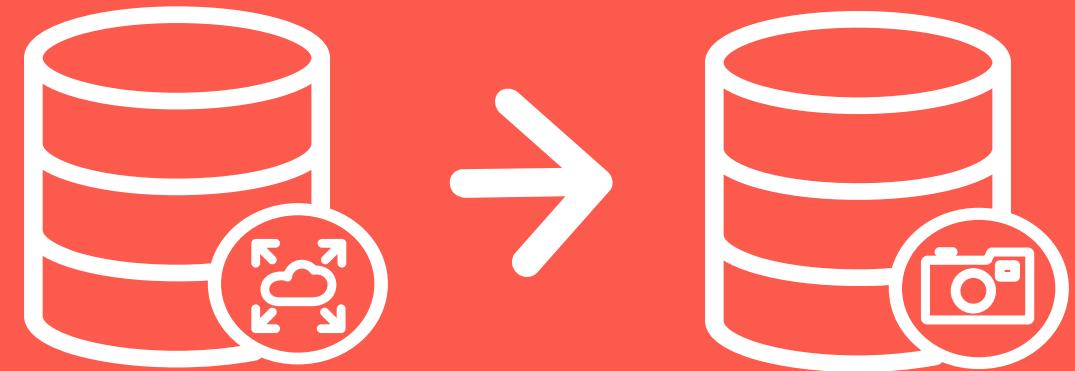


Ease of use

usually offers a bunch of prebuilt connectors that accelerate integration projects

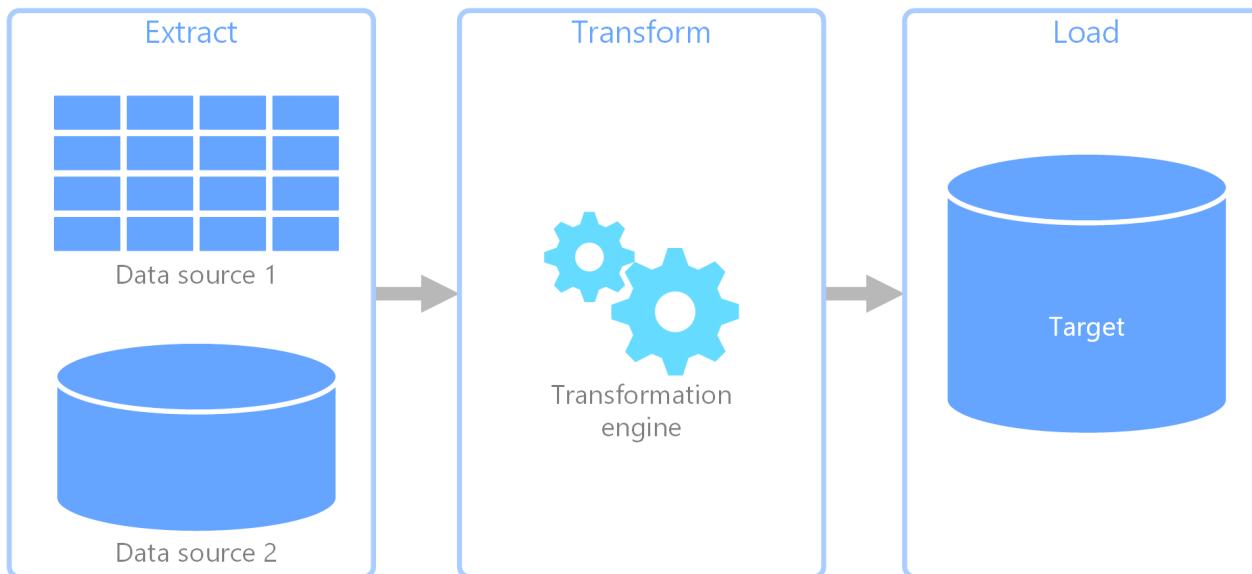
Consideration: Snapshots

- ✓ Since migrations are theoretically executed at a point in time, any data which is created post execution will not be brought over
- ✓ To synchronize the *delta* you can use the broadcast pattern (which we'll see next)
- ✓ Keep in mind that data can still be created during the time that the migration is running unless you prevent it or create a separate table where new data can be stored



Consideration: Transforming the Dataset

Other than a few exceptions, like data backup/restore or migrating to new hardware, you usually need to modify the formats and structures of data during a migration



Extract

Process of reading data from a database

Transform

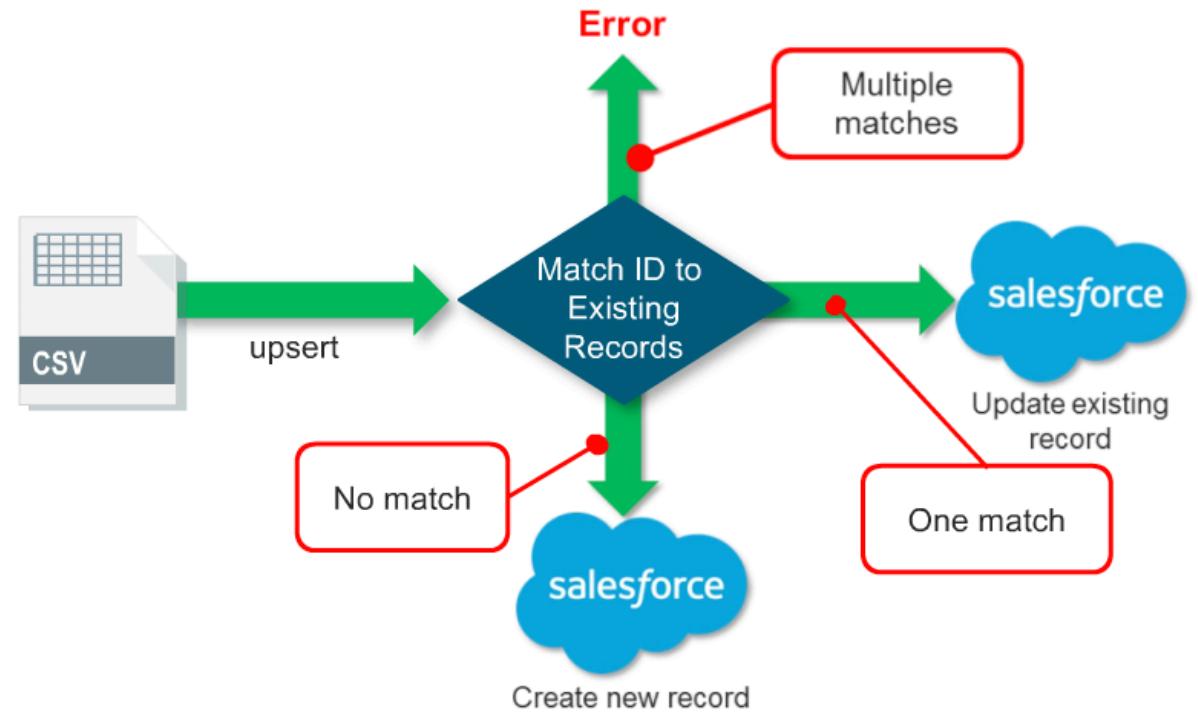
Process of converting extracted data from its previous form into the form it needs to be in so that it can be placed into another database

Load

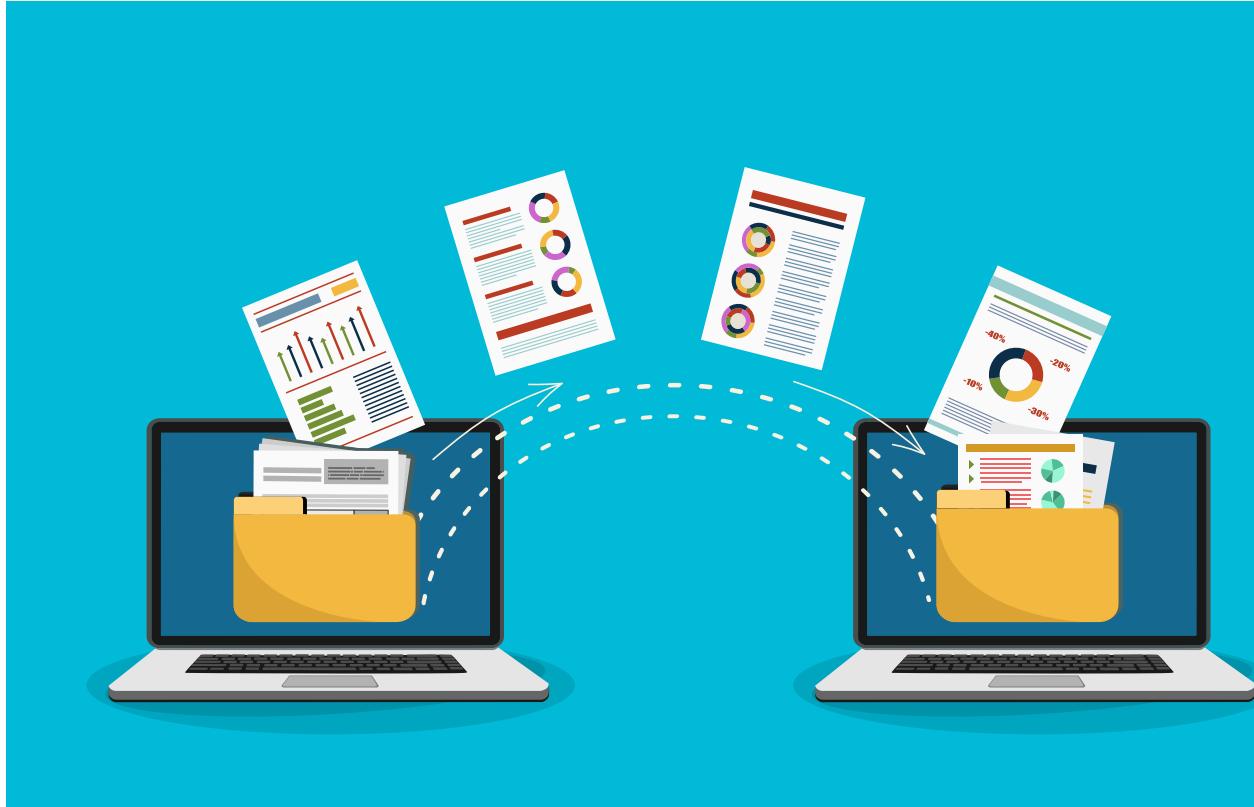
Process of writing the data into the target database

Consideration: Insert vs Update

- ✓ When running a migration multiple times you may need to handle both creating and updating records
- ✓ Duplicate or incomplete entries can be extremely costly
- ✓ Intelligently picking the right fields to **upsert** on is especially important if the target system is the system of record
- ✓ Having a solution that checks for duplicates can make a huge difference, both proactively and reactively



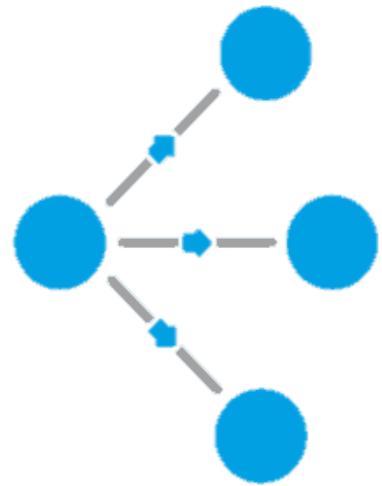
Summary



Migration is the one time act of taking a specific data set from one system and moving it to another



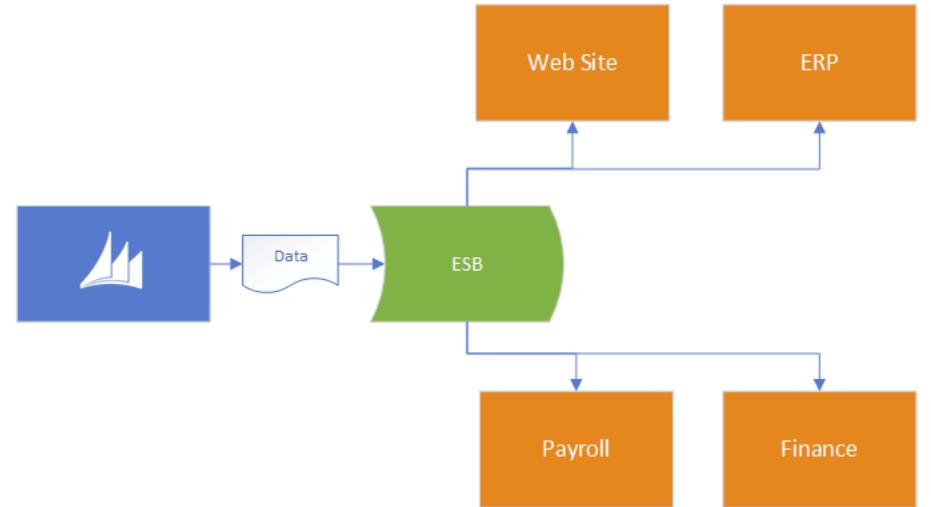
Creating an application for the purpose of migrations is valuable because you can make it a reusable service



Broadcast
one way sync from one to many

What is It?

-  Move of data from a single source system to many destinations
-  Keeps data up to date between multiple systems, across time
-  Like the migration pattern, only moves data in one direction, from the source to the destination



Migration vs. Broadcast

What separates using the broadcast pattern and the migration pattern set to automatically run every few seconds?

Migration

- Full scope
- large volumes
- Parallel processing
- graceful failure

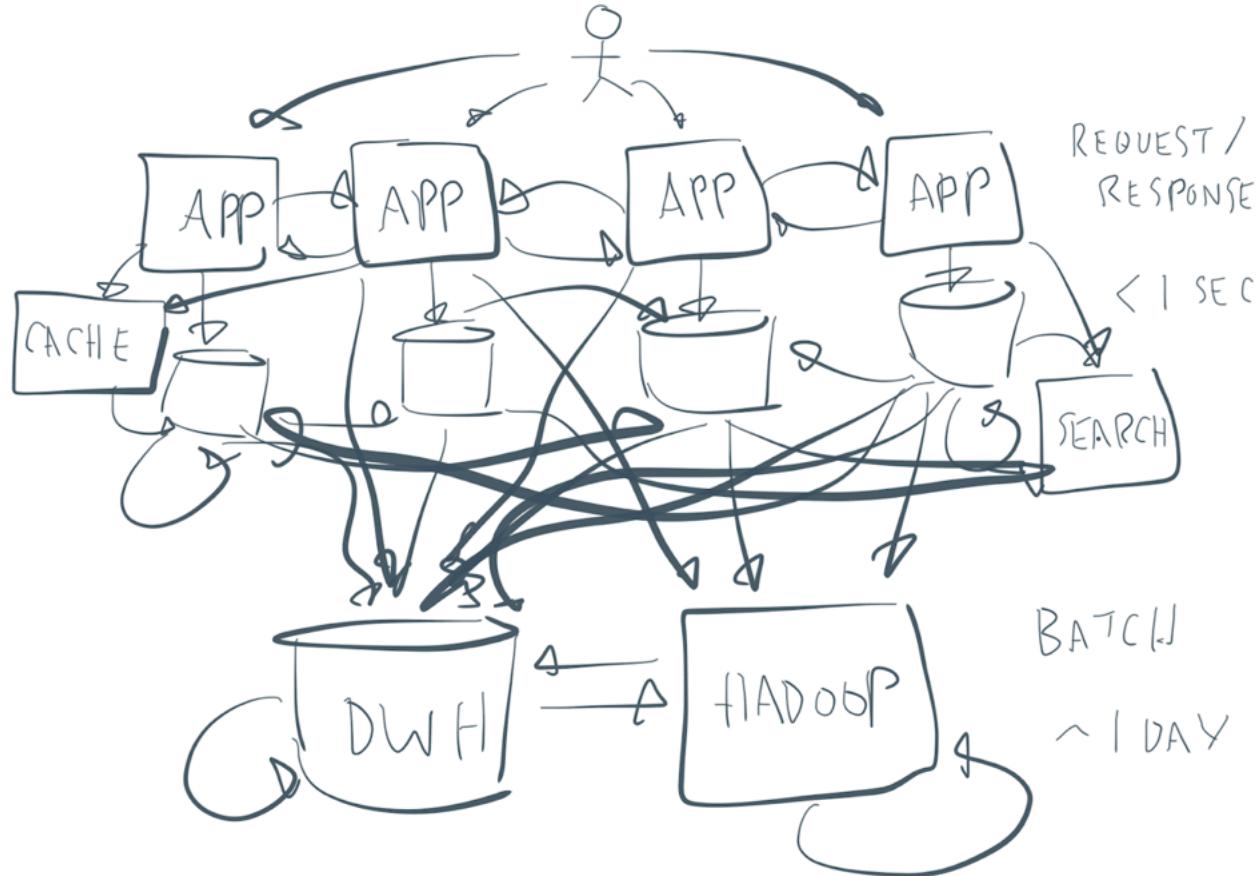
Broadcast

- Change capture
- “as quickly as possible”
- transactional
- highly reliable and available

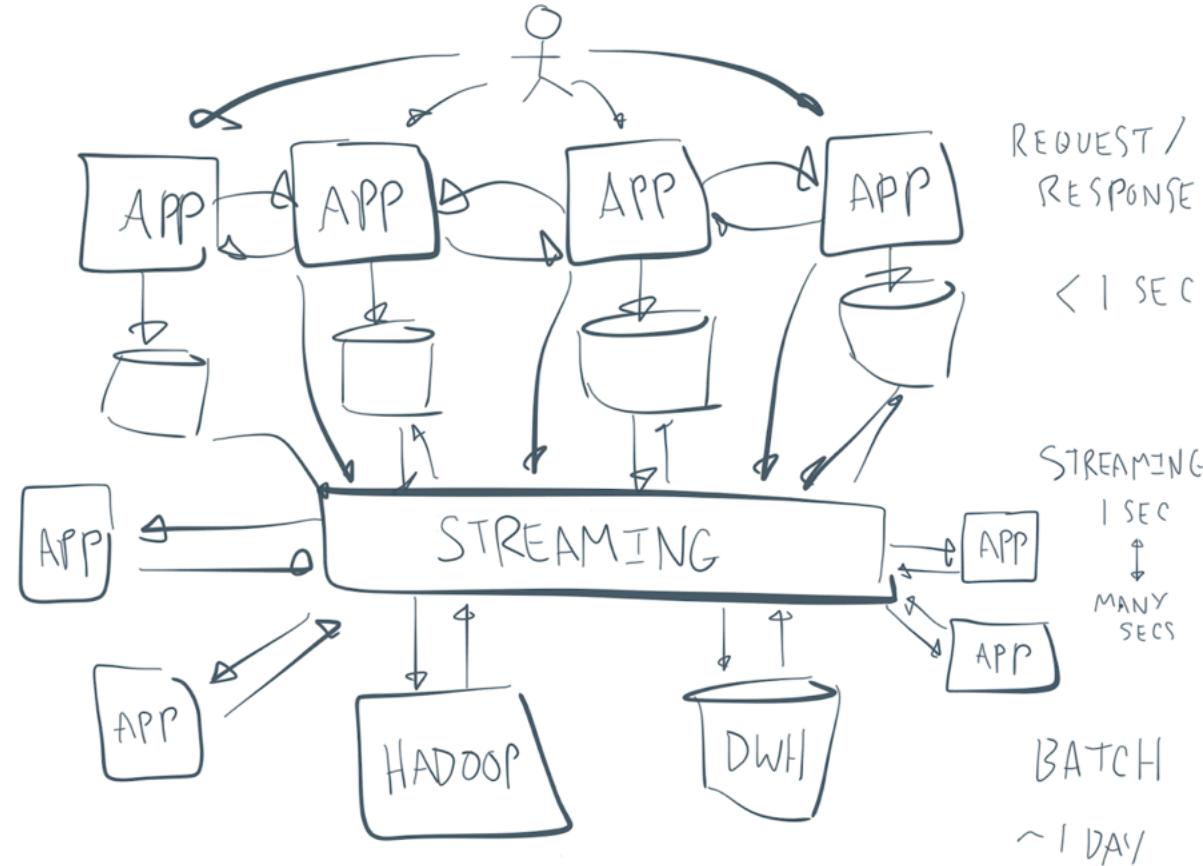
Transactional

If a data transfer (“transaction”) succeeds, data is committed (i.e. persisted) at destination

Note about EDA (Event Driven Architecture)



Note about EDA (Event Driven Architecture)



Simply put, the event is a significant change in state, which is triggered when a user takes an action.

Why is it valuable?

Handles any generic situation where a system needs to know information in near real time residing in another system



Use Case I

A real time reporting dashboard that receives updates from multiple systems in real time



Use Case II

A fulfilment processing system that immediately starts processing orders from your CRM, online e-shop, or internal tool



BROADCAST PATTERN

When is it Useful?



A system needs to know as soon as an event in another system happens



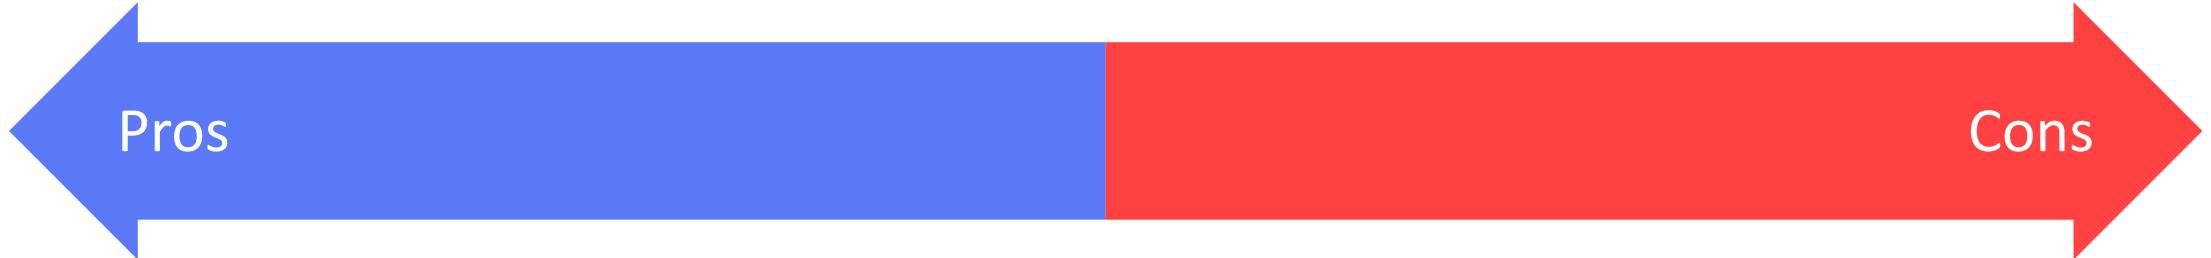
Data needs to flow from A to B automatically, without human involvement



System A doesn't need to know what happens with the object in system B

Consideration: Input Flow

Option 1: Data Passed via Message Notification Payload



Pros

minimizing the number
of API calls



Cons

code and/or configuration that sends notification
needs to be developed differently for each source



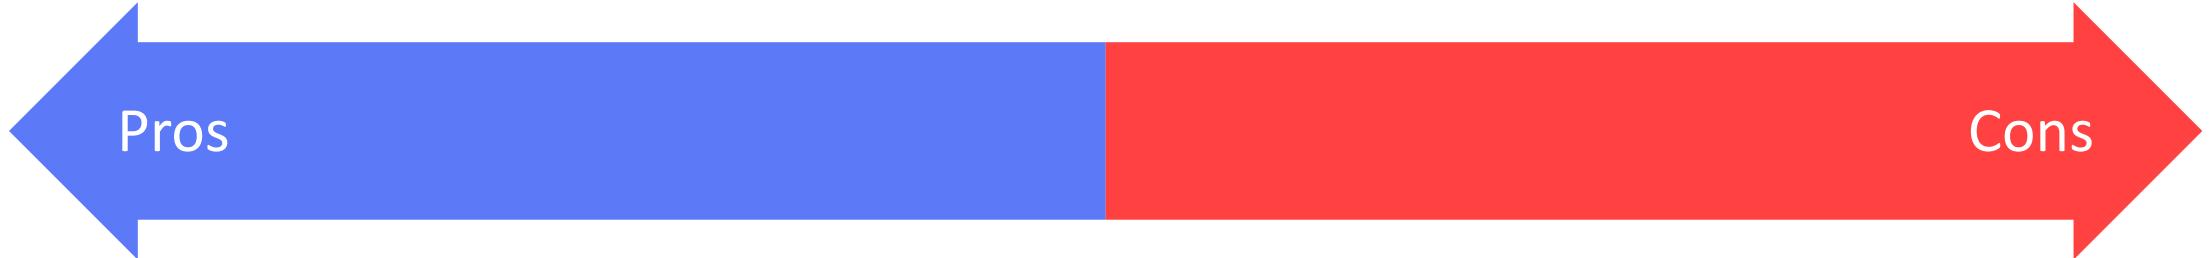
can become “chatty” when the source
system has a lot of changes



if the integration application is not
running properly or is down you could
lose data

Consideration: Input Flow

Option 2: Integration application pulls data from the source system



Pros

pull will only run when the application is healthy



watermark only updated once message is successfully processed



Cons

still requires the development of the outbound message notifying the integration application that there is data waiting to be pulled



can become “chatty” when the source system has a lot of changes



Consideration: Object Payload

-  Moving field values as payloads is very expensive in terms of development cost and performance
-  This means synchronizing whole objects even though only one field has changed, exposing a problem when you have multiple broadcasts feeding into the same system

```
{  
  "id": "1",  
  "type": "cars",  
  "cars_array": [  
    {  
      "id": "1",  
      "name": "Mercedes Benz",  
      "desc": "Mercedes Benz is a global automobile marque and a division of the German company Daimler AG. The brand is known for luxury vehicles, buses, coaches, and lorries. The headquarters is in Stuttgart, Baden-Württemberg. The name first appeared in 1926 under Daimler-Benz."  
    },  
    {  
      "id": "2",  
      "name": "BMW",  
      "desc": "BMW AG, originally an initialism for Bayerische Motoren Werke in German, or Bavarian Motor Works in English) is a German multinational company which currently produces luxury automobiles and motorcycles, and also produced aircraft engines until 1945."  
    }  
  ]  
}
```

Consideration: Field Data

SCENARIO

CRM A and CRM B both broadcasting contacts into CRM C and you need to handle cases where both systems change the same or different fields on the same contact

Option 1

Merges the values of the two payloads, with the later one winning if the same field are affected

Option 2

Define a master in the integration application configuration at either the system, object or field level



Exception Handling

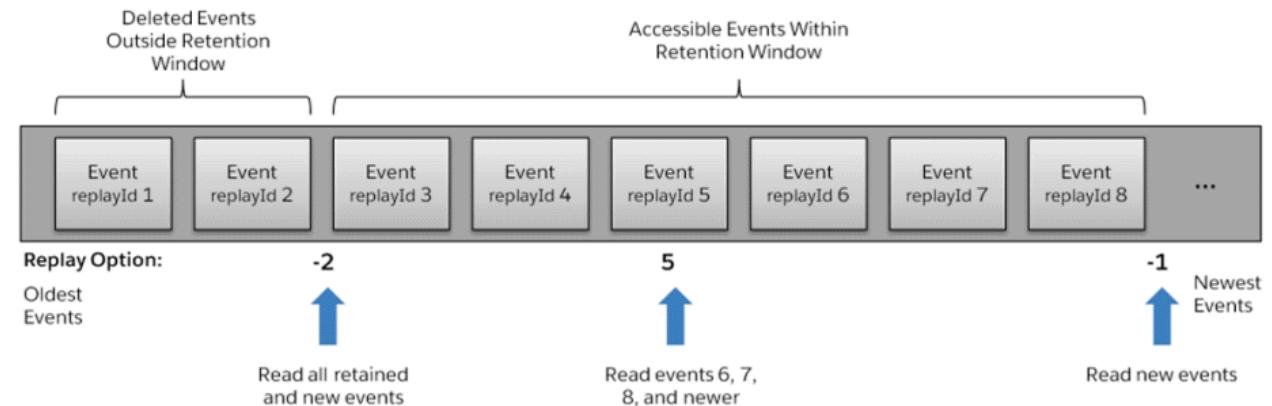
- A simple exception management strategy can be to log all the issues to a file or send in an email
- Create a notification workflow based on the values in the payload so that different people are notified based on different errors for those exceptions that cannot be fixed by a prescriptive design time solution.

```
try {
    something();
}

catch(e) {
    window.location.href =
    "http://stackoverflow.com/search?q=[js]+" + e.message;
}
```

Consideration: Watermark (Replay)

- Keeps track of when the last poll was done such that we grab only the newest set of objects that have been modified in that time window
- In systems without watermarking, you may have to either process all the records, or produce a clever way to capture only those items to be processed
- In general, using a *last modified timestamp* as the replayid provides the cleanest way to pull only the latest content



Summary



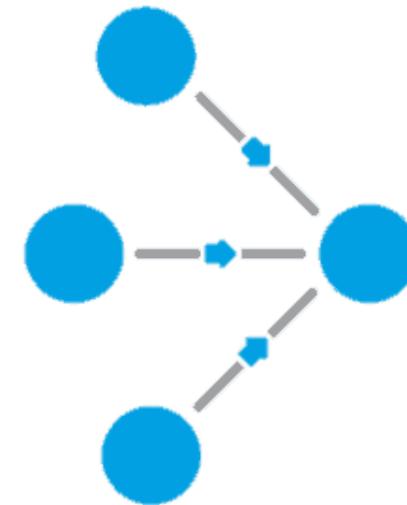
Defined as the act of moving data from one to one or more systems in as near real-time fashion



Valuable when you have data in one system and would like to have that data in other systems as soon as possible.

Aggregation

The act of taking or receiving data from multiple systems and inserting into one



SCENARIO

Your company has customer data in three different systems, and you need to generate a report which uses data from all three of those systems.

OPTION 1

a daily migration from each of those systems to a data repository and then query against that database.

- ✓ would have a whole other database to worry about and keep synchronized.
- ✓ As things change in the three other systems, would have to constantly ensure that the data repository is kept up to date
- ✓ Information would be a day old

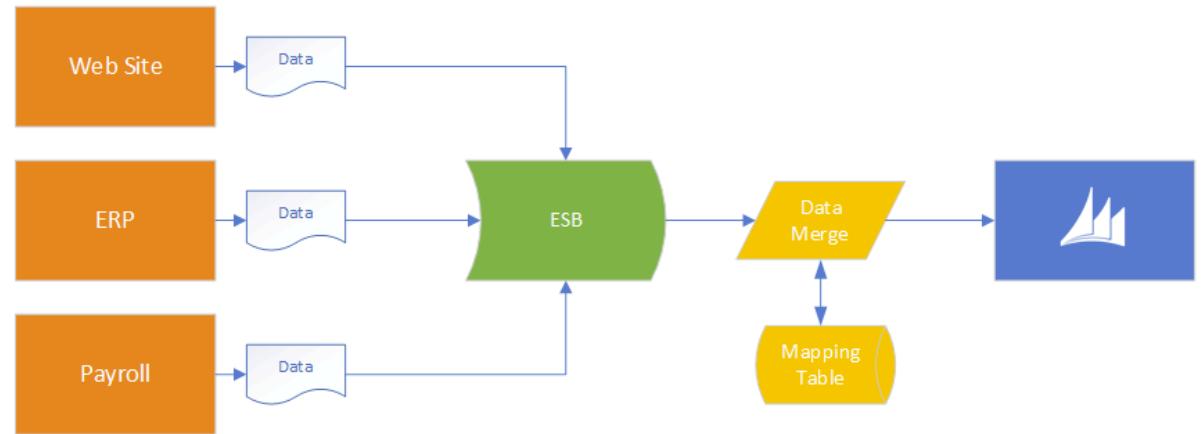
OPTION 2

set up three broadcast applications so that the reporting database is always up to date

- ✓ would need to maintain a database whose only purpose is to store replicated data so that I can query it every so often.
- ✓ Not to mention the number of wasted API calls to ensure that the database is always up to x minutes from reality.

What is It?

- ✓ The act of taking or receiving data from multiple systems and inserting into one
- ✓ On demand query multiple systems merge the data set, and do as you please with it
- ✓ From our scenario, you can build an integration app which queries the various systems, merges the data and then produces a report



Why is it valuable?

- ✓ Allows you to extract and process data from multiple systems in one application
- ✓ Data is up to date at the time that you need it, does not get replicated, and can be processed/merged to produce the dataset you want

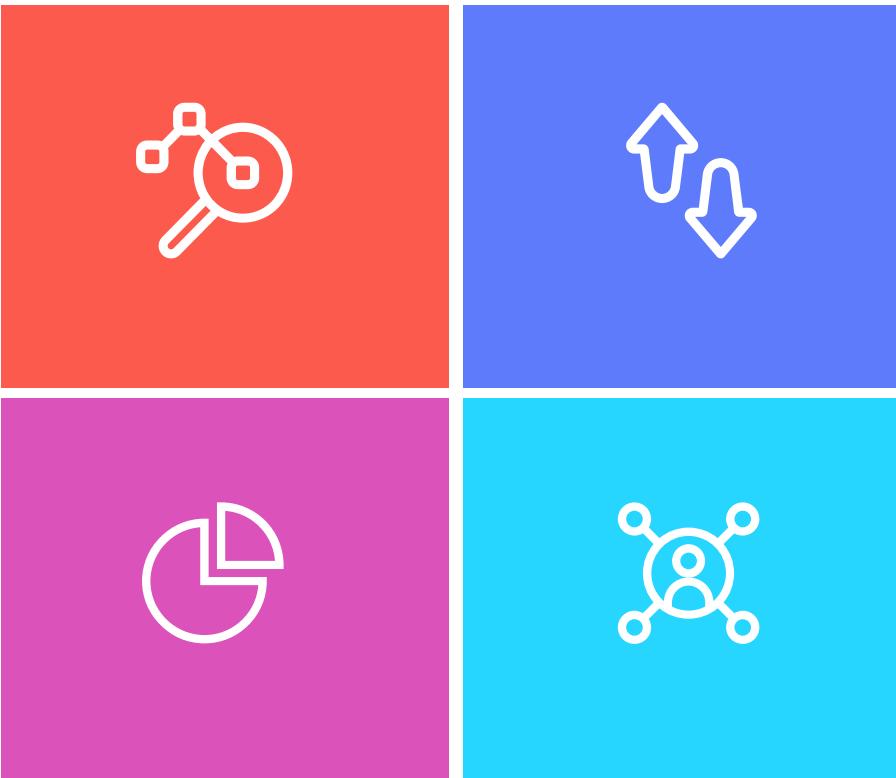


When is it useful?

The aggregation pattern is valuable if you are creating:

Orchestration APIs

to “modernize” legacy systems



Reports or Dashboards

which must pull data from multiple systems and create an experience with that data

Compliance

auditing systems which need to have related data from multiple systems

Ease of Use

reduce the amount of learning that needs to take place across the various systems to ensure visibility into what is going on

Consideration: Collecting Data



Listener

create a listener which waits for messages from multiple systems and aggregates them in real time



Trigger

create an application that is triggered, processes, and returns the results

In general, reserve the aggregation pattern for collection of data on demand or for orchestration of data rather than synchronization

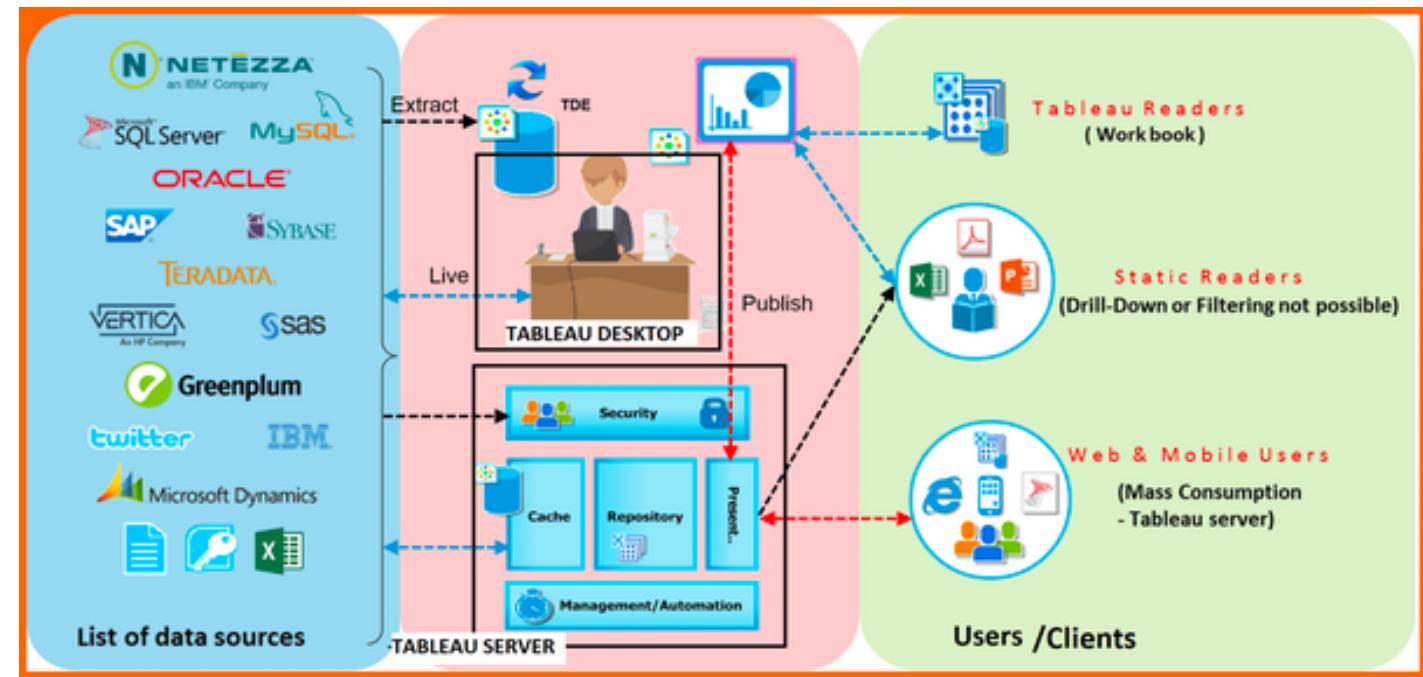
Consideration: Migration from Many to One vs Aggregation



Although you can create a migration that collects data from many systems and then inserts them into one system, a migration is realistically the movement of existing data to the destination system where the data will reside.



aggregation is meant to correlate, merge, and potentially transform the object and usually not storing the new data set into a system where it will continue to be used while deprecating the origin.



Consideration: Source Data Scope

✓ Since we use different connectors for each system we are pulling data from you have the ability to write different queries to pull data from those systems

✓ This means that you can do things like get accounts that are older than 2 years in one system and account that are older than 1 year in the other



Consideration: Merging the multiple datasets

- Once you have queried across the systems, you will most likely have stored the results in the local store for the duration of the app running. At this point, you need to choose how to merge the datasets or what to do with it.



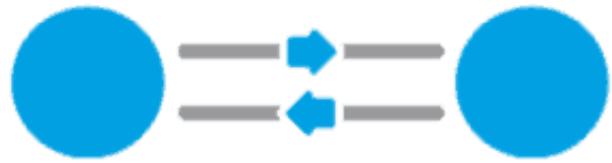
Summary



The aggregation pattern is very useful when building APIs that collect and return data from multiple systems



can be modified to merge and format the data as you would like and can be easily extended to insert the data into multiple systems



Bi-directional Synchronization

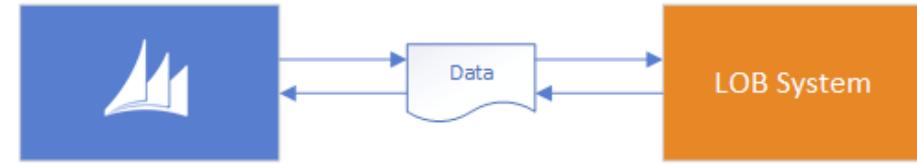
/baɪdɪ'rekʃənəl/

adjective

Moving or operating in two usually opposite directions

What is It?

-  The act of unioning two datasets in two different systems to behave as one while respecting their need to exist as different datasets
-  Different tools or different systems for accomplishing different functions on the same data set



BI-DIRECTIONAL SYNC PATTERN

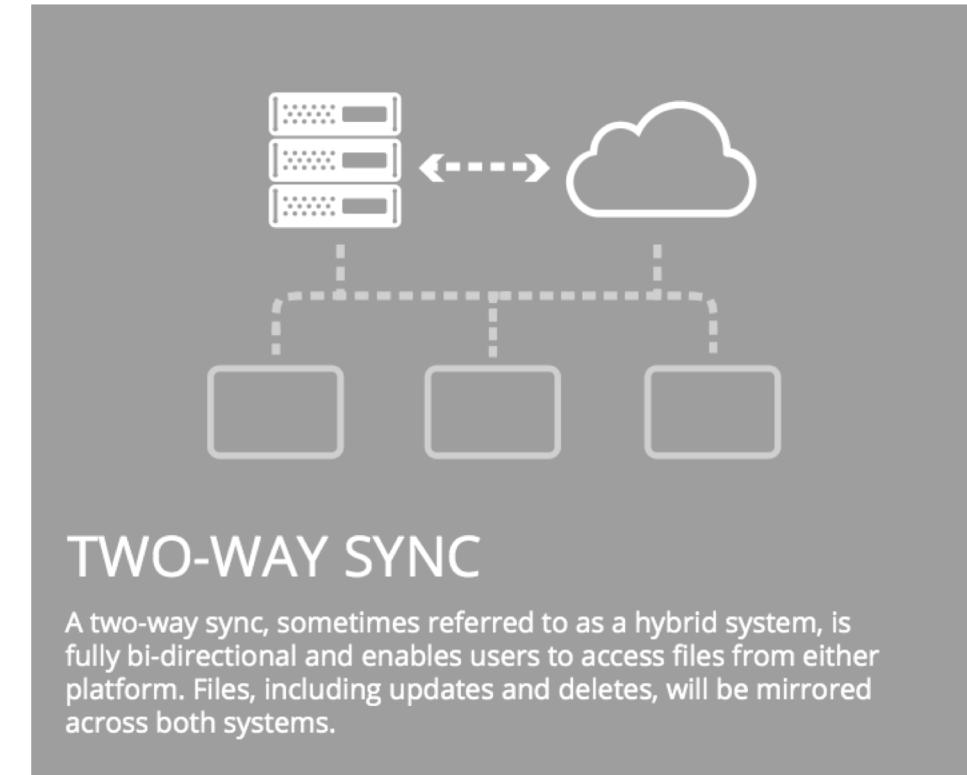
Why is it valuable?



In a situation where you have two or more independent, isolated, representations of the same reality



Bi-directional sync takes you from a suite of products that work well together but may not be the best at their own individual function, to a suite that you hand pick and integrate together using an integration platform.



Why is it valuable?

Single View of the Customer

Manually giving everyone access to all the systems that have a representation of the notion of a customer can hinder productivity, training, security, and cost

Bi-directional synchronization allows users to have a realtime view of the same customer within the perspective that they care about.



Your Data Expectations vs Reality

Consideration: Filtering and Object Mapping

SCENARIO

You have two Salesforce instances, one for the Ohio Sales team, and one for the rest of the US. If you implemented a bidirectional sync, instead of always moving all contacts from one to the other, you could set a filter like “From OH>US move all contacts that are in US but not OH” and “From US->OH move all contacts that are OH”.



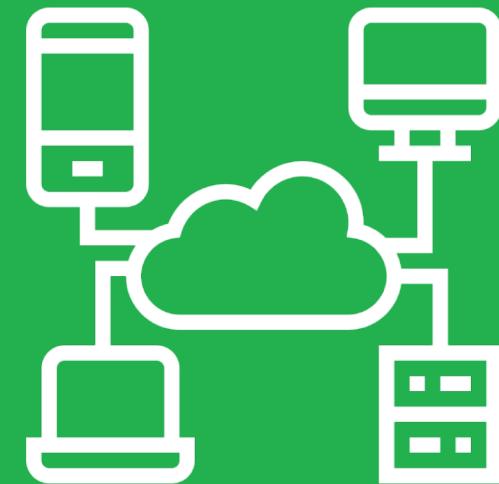
Like a pair of broadcast application with the ability to scope and map the data in both directions



With these kinds of filters you will make it so that you don't synchronize contacts which make sense in your instance, but have no value in the other instance

Consideration: 3 or More Systems

- Need to change to a different architecture and start introducing the notion of master
- Consider a hub and spoke architecture:
one system will be the system of record for information needing synchronized
- Alternatively. decouple the systems using queues and a pub-sub model where every system will broadcast its change to a queue and system can look for changes that they care about





Summary



Correlation

/kôrə'lāSH(ə)n/

noun

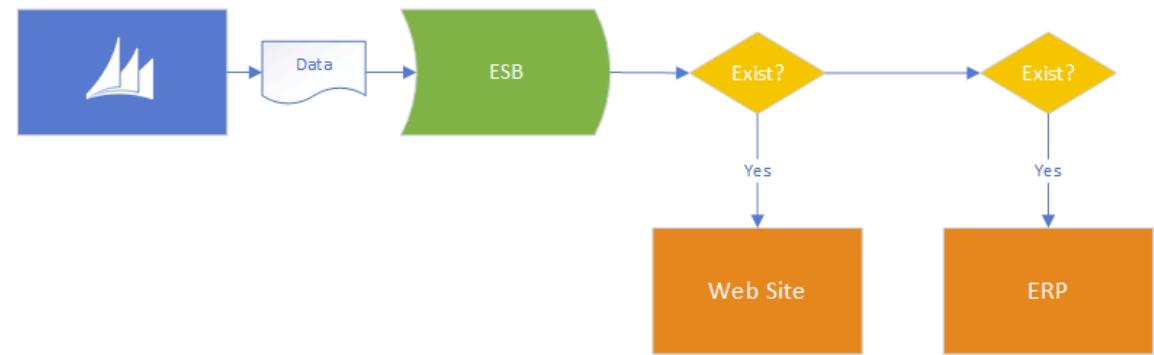
1. a mutual relationship or connection between two or more things.

"research showed a clear **correlation** between recession and levels of property crime"



What is It?

- Identifies the intersection of two data sets and does a bi-directional synchronization only if record occurs in both systems naturally
- Unlike the bi-directional pattern, which synchronizes the union of the scoped dataset, correlation synchronizes the intersection
- Does not care where objects came from; agnostically synchronize as long as data is found in both systems



Why is it valuable?

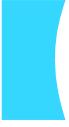


Useful in the case where you have two groups or systems that want to share data only if they both have a record representing the same item/person in reality



Only bi-directionally synchronizes the objects on a “Need to know” basis rather than always moving the full scope of a dataset in both directions

“Hey, lets share data between the two hospital so if a patient uses either hospital, we will have a up to date record of what treatment they received at both hospitals.”



CORRELATION PATTERN

When is useful?

SCENARIO



Consideration: Bi-Directional Factor

What are key things to keep in mind when building applications using the correlation pattern?

Given that the correlation pattern is almost the same as the bi-directional sync, the same considerations apply.

The main difference is in the impact of how you define the Same.

Consideration: Definition of Same

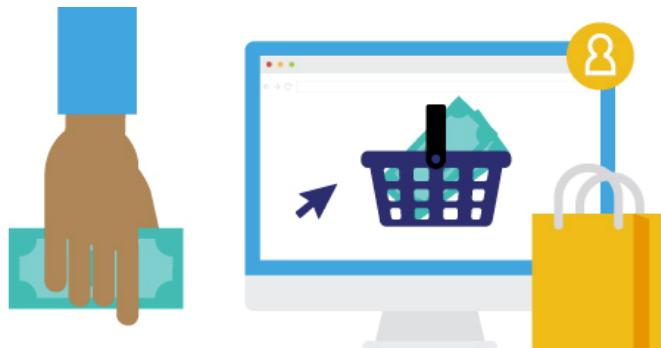
Definition differs by industry and use case and the degree to which it is harmful will also differ by industry and use case

	Just Right	Too Strict	Too Loose
Bi-Directional Sync	Union of Data Sets	Duplicate Created	Wrong Records Merged
Correlation	Intersection of Data Sets	Not Synchronized	Wrong Records Merged

Consideration: Definition of Same

When keeping it real goes wrong

What would happen if the criteria that evaluates if two records in two different systems are the same is incorrectly defined?



SCENARIO

An ecommerce site where the definition of “same customer” is defined by a composition of a customer’s first and last name

If we send recommendations via email, We would end up with an error where two people with the same name would receive the union of both sets recommendations

Consideration: Definition of Same

SCENARIO

Two babies named Kylo are born in the same hospital, with the last name Ren. Patient uniqueness in the hospital's systems is determined by

SSN + First Name + Last Name + City + Hospital



It is paramount that you are very careful when picking the right *definition of unique* when dealing with high sensitivity information



Likely they will share everything except SSN, so they likely are only apart in the hash by 1 wrong key stroke when filling in their info

Summary



"It's important to remember that correlation does not imply causation. Besides, we all know it was Brian."



The intersection of datasets as opposed to bi-directional sync which is the union of the datasets



Allows you to keep two systems in sync at a lower "cost"

Wrap-up

- ✓ Following these integration patterns will help the scalability and maintainability of your integrations
- ✓ Each of these patterns can be applied by themselves or composed
- ✓ Patterns have a time and place; “do what makes sense”



Thank You!